Continuous Random Variables

James Heald[1]

[1]Gatsby Computational Neuroscience Unit
University College London

Gatsby Bridging Programme 2023

Objectives

- Introduce the concept and formal definition of a continuous random variable X and a probability density function.
- Learn how to find the probability that a continuous random variable falls in some interval [a, b].
- Learn that if X is continuous, the probability that X takes on any specific value is 0.
- Introduce the concept and formal definition of a cumulative distribution function of a continuous random variable.
- Learn how to find the cumulative distribution function of a continuous random variable X from the probability density function of X.

Discrete vs. continuous random variables

Unlike discrete random variables, which can take on a countable number of possible values (e.g. faces of a die or cards in a deck), continuous random variables can take on an uncountable number of possible values (e.g. all the real numbers in some interval).

Discrete vs. continuous random variables

Unlike discrete random variables, which can take on a countable number of possible values (e.g. faces of a die or cards in a deck), continuous random variables can take on an uncountable number of possible values (e.g. all the real numbers in some interval).

Examples

- the voltage membrane potential of a cell
- the interspike interval of a neuron
- the force generated by a muscle
- the velocity of an eye movement

Continuous random variables

Definition

A random variable X is continuous if:

1. possible values comprise either a single interval on the number line (i.e. for some $a < b$, any number x between a and b is a possible value) or a union of disjoint intervals, and
2. $P(X = c) = 0$ for any number c that is a possible value of X.

Discrete probability distributions in the limit

Continuous random variables can be discretised into bins to form a probability mass function. As the bins become narrower, the probability mass function approaches a smooth curve.

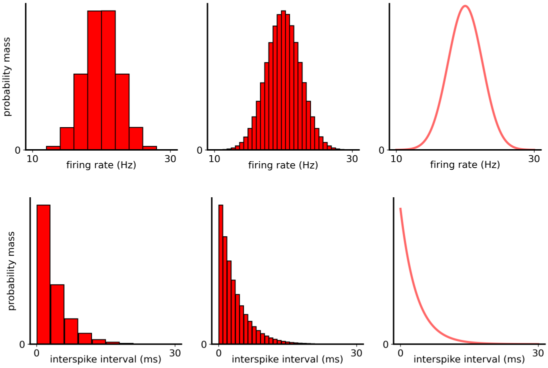

Probabilities as integrals

The probability that a continuous random variable X takes on a value in the interval [a, b] is given by the area under the probability density function $f(x)$.

$$P(a \leq X \leq b) = \int_a^b f(x)dx$$

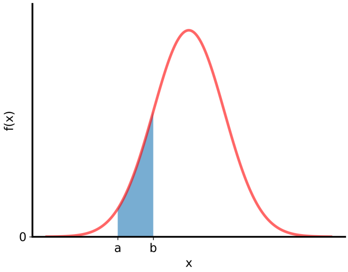

The probability density function

Definition

A random variable X is continuous if there exists a nonnegative function $f(x)$ defined on the interval $(-\infty, \infty)$, such that for any interval $[a, b]$ we have

$$P(a \leq X \leq b) = \int_a^b f(x)dx.$$

The probability density function

Definition

A random variable X is continuous if there exists a nonnegative function $f(x)$ defined on the interval $(-\infty, \infty)$, such that for any interval $[a, b]$ we have

$$P(a \leq X \leq b) = \int_a^b f(x)dx.$$

A valid probability density function (PDF) $f(x)$ has the following properties:

$$f(x) \geq 0 \text{ for all } x \tag{1}$$

$$\int_{-\infty}^{\infty} f(x)dx = 1. \tag{2}$$

Density as probability per unit length

The probability of a small interval δ is approximately the density $\times\ \delta$:

$$\begin{aligned} P(x \leq X \leq x+\delta) &= \int_x^{x+\delta} f(t)dt \\ &\approx f(x) \times \delta \end{aligned}$$

Thus density is probability per unit length (probability accumulation rate):

$$\frac{P(x \leq X \leq x+\delta)}{\delta} \approx f(x)$$

Each possible value has zero probability

The probability that X takes on a particular value a is 0, as

$$\begin{aligned} P(X = a) &= \int_a^a f(x)dx \\ &= \lim_{\epsilon \to 0} \int_{a-\epsilon}^{a+\epsilon} f(x)dx \\ &= 0. \end{aligned}$$

This implies that probabilities don't depend on interval end points:

$$P(a \le X \le b) = P(a < X < b) = P(a < X \le b) = P(a \le X < b),$$

as $P(X = a) = P(X = b) = 0$.

Mean and variance

The expected value (mean) of a continuous random variable X is:

$$\mathbb{E}[X] = \int_{-\infty}^{\infty} x f(x) \mathrm{d}x,$$

the expected value of a function $g(x)$ of X is:

$$\mathbb{E}[g(X)] = \int_{-\infty}^{\infty} g(x) f(x) \mathrm{d}x,$$

the variance of X is:

$$\begin{aligned} \mathrm{Var}[X] &= \mathbb{E}[(x - \mathbb{E}[X])^2] \\ &= \int_{-\infty}^{\infty} (x - \mathbb{E}[X])^2 f(x) \mathrm{d}x. \end{aligned}$$

Example: the uniform distribution

When X has a uniform distribution, the PDF is:

$$f(x) = \begin{cases} \frac{1}{b-a} & \text{if } a \leq x \leq b \\ 0 & \text{otherwise,} \end{cases}$$

the expected value of X is:

$$\mathbb{E}[X] = \int_a^b x \frac{1}{b-a} \mathrm{d}x = \frac{a+b}{2},$$

the variance of X is:

$$\mathrm{Var}[X] = \int_a^b \left(x - \frac{a+b}{2}\right)^2 \frac{1}{b-a} \mathrm{d}x = \frac{(b-a)^2}{12}.$$

The cumulative distribution function (CDF)

The cumulative distribution function $F(x)$ is the area under the probability density function $f(x)$ to the left of x.

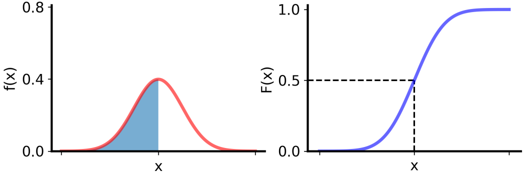

The cumulative distribution function (CDF)

Definition

Let X be a continuous random variable with probability density function $f(x)$, then the cumulative distribution function is defined as

$$\begin{aligned} F(x) &= P(X \leq x) \\ &= \int_{-\infty}^{x} f(t)dt. \end{aligned}$$

The CDF is a monotonically-increasing continuous function $F : \mathbb{R} \mapsto [0, 1]$ satisfying $\lim_{x \to -\infty} F(x) = 0$ and $\lim_{x \to \infty} F(x) = 1$.

Computing probabilities using the CDF

$$P(a \leq X \leq b) = F(b) - F(a).$$

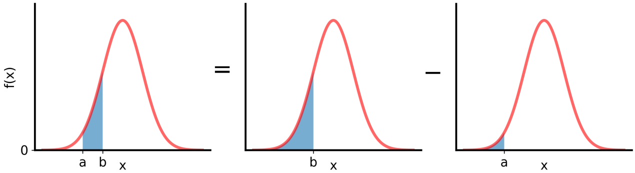

$$P(X > a) = F(\infty) - F(a)$$
$$= 1 - F(a).$$

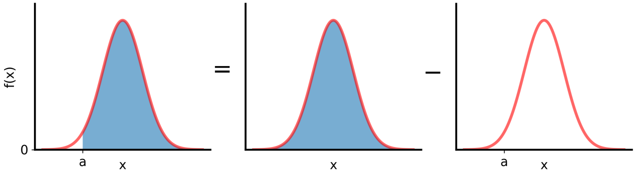

Obtaining the PDF from the CDF

At every x at which the derivative $F'(x)$ exists, $F'(x) = f(x)$.

Examples

When X has a uniform distribution, for $a < x < b$:

$$F'(x) = \frac{d}{dx}\left(\frac{x-a}{b-a}\right) = \frac{1}{b-a} = f(x)$$

Sampling using the CDF

The inverse transform sampling algorithm can be used to sample a continuous random variable using the inverse of its cumulative distribution function.

Recall that $F : \mathbb{R} \mapsto [0, 1]$.

To draw a sample $x \sim f(x)$:

1. Sample $u \sim \mathrm{U}(0, 1)$ (using a pseudo-random number generator)
2. Let $x = F^{-1}(u)$

Example: the exponential distribution

The PDF of the exponential distribution is:

$$f(x) = \begin{cases} \lambda e^{-\lambda x} & \text{if } x \geq 0 \\ 0 & \text{otherwise,} \end{cases}$$

which implies that the CDF is:

$$F(x) = 1 - e^{-\lambda x} = u,$$

and the inverse of the CDF is:

$$F^{-1}(u) = -\frac{\log(1-u)}{\lambda} = x.$$

Hence, to sample $x \sim f(x)$:

1. Sample $u \sim \mathrm{U}(0, 1)$
2. Let $x = -\frac{\log(1-u)}{\lambda}$

The normal (Gaussian) distribution