The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Inference

Joaquín Rapela

Gatsby Computational Neuroscience Unit
University College London

July 18, 2023

Contents

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1. The Gaussian distribution
2. Linear models for regression
   - Least-squares regression
   - Maximum-likelihood regression
   - Bayesian linear regression
     - Batch Bayesian linear regression
     - Online Bayesian linear regression

Main reference

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

I will mainly follow chapters two *Probability distributions* and three *Linear models for regression* from Bishop (2016).

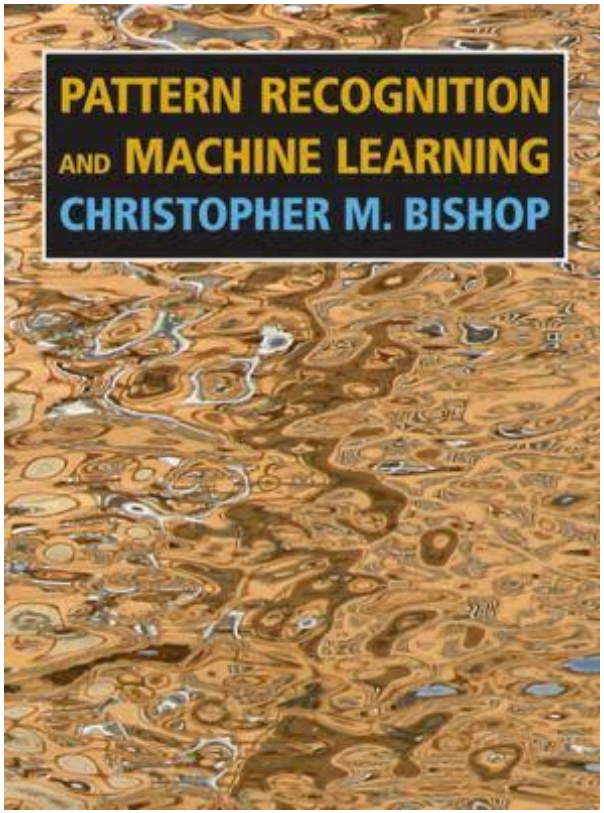

Contents

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1. The Gaussian distribution
2. Linear models for regression
   - Least-squares regression
   - Maximum-likelihood regression
   - Bayesian linear regression
     - Batch Bayesian linear regression
     - Online Bayesian linear regression

The Gaussian distribution

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

- One-dimensional

$$\mathcal{N}(x|\mu,\sigma^2) = \frac{1}{(2\pi)^{\frac{1}{2}}(\sigma^2)^{\frac{1}{2}}} \exp\left\{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}\right\}$$

- D-dimensional

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu},\Sigma) = \frac{1}{(2\pi)^{D/2}\Sigma^{\frac{1}{2}}} \exp\left\{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^{\mathsf{T}}\Sigma^{-1}(\mathbf{x}-\boldsymbol{\mu})\right\}$$

The Gaussian is the maximum entropy distribution (Cover and Thomas, 1991)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Definition 1 (Differential entropy)

The differential entropy $h(X)$ of a continuous random variable X with a density $f(x)$ is defined as

$$h(X) = -\int_S f(X) \log f(x)\ dx$$

where S is the support set of the random variable.

Theorem 1 (The Gaussian is the maximum entropy distribution)

Let the random vector $X \in \mathbb{R}^n$ have zero mean and covariance K. Then $h(X) \leq \frac{1}{2} \log(2\pi e)^n |K|$, with equality if $X \sim \mathcal{N}(0, K)$.

The central limit theorem (Papoulis and Pillai, 2002)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Theorem 2 (The central limit theorem)

Given n independent and identically distributed random vectors $\mathbf{X}_i$, with mean vector $\boldsymbol{\mu} = E\{\mathbf{X}_i\}$ and covariance matrix $\boldsymbol{\Sigma}$. Then

$$\sqrt{n}(\bar{\mathbf{X}}_n - \boldsymbol{\mu}) \to \mathcal{N}(0, \Sigma)$$

with convergence in distribution.

Very useful properties of the Gaussian distribution (Bishop, 2016)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Theorem 3 (Marginals and conditionals of Gaussians are Gaussians)

Given $\mathbf{x} = \begin{bmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{bmatrix}$ *such that*

$$p(\mathbf{x}) = \mathcal{N}\left(\mathbf{x} \middle| \begin{bmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{bmatrix}, \begin{bmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{bmatrix}\right)$$
$$= \mathcal{N}\left(\mathbf{x} \middle| \begin{bmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{bmatrix}, \begin{bmatrix} \Lambda_{aa} & \Lambda_{ab} \\ \Lambda_{ba} & \Lambda_{bb} \end{bmatrix}^{-1}\right)$$

Then

$$p(\mathbf{x}_a|\mathbf{x}_b) = \mathcal{N}\left(\mathbf{x}_a \middle| \boldsymbol{\mu}_a - \Lambda_{aa}^{-1}\Lambda_{ab}(\mathbf{x}_b - \boldsymbol{\mu}_b), \Lambda_{aa}^{-1}\right) \quad (1)$$
$$= \mathcal{N}\left(\mathbf{x}_a \middle| \boldsymbol{\mu}_a + \Sigma_{ab}\Sigma_{bb}^{-1}(\mathbf{x}_b - \boldsymbol{\mu}_b), \Sigma_{aa} - \Sigma_{ab}\Sigma_{bb}^{-1}\Sigma_{ba}\right) \quad (2)$$
$$p(\mathbf{x}_b) = \mathcal{N}\left(\mathbf{x}_b \middle| \boldsymbol{\mu}_b, \Sigma_{bb}\right) \quad (3)$$

Very useful properties of the Gaussian distribution (Bishop, 2016)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Theorem 4 (Marginals and conditionals of the linear Gaussian model)

Given the linear Gaussian model

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \Lambda^{-1})$$
$$p(\mathbf{t}|\mathbf{x}) = \mathcal{N}(\mathbf{t}|A\boldsymbol{\mu} + \mathbf{b}, L^{-1})$$

Then

$$p(\mathbf{t}) = \mathcal{N}(\mathbf{t}|A\boldsymbol{\mu} + \mathbf{b}, L^{-1} + A\Lambda^{-1}A^{\intercal}) \tag{4}$$
$$p(\mathbf{x}|\mathbf{t}) = \mathcal{N}(\mathbf{x}|\Sigma\{A^{\intercal}L(\mathbf{t} - \mathbf{b}) + \Sigma\boldsymbol{\mu}\}, \Sigma) \tag{5}$$

where

$$\Sigma = (\Lambda + A^{\intercal}LA)^{-1}$$

Very useful properties of the Gaussian distribution (Bishop, 2016)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

The conditional, $p(\mathbf{x}|\mathbf{t})$, of the linear Gaussian model is the fundamental result used in the derivation of

1. Bayesian linear regression (Bishop, 2016),
2. Gaussian process regression (Williams and Rasmussen, 2006),
3. Gaussian process factor analysis (Yu et al., 2009),
4. linear dynamical systems (Durbin and Koopman, 2012).

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Claim 1 (Quadratic form of Gaussian log pdf)

$p(\mathbf{x})$ is a Gaussian pdf with mean $\boldsymbol{\mu}$ and precision matrix Λ if and only if $\int p(\mathbf{x})d\mathbf{x} = 1$ and

$$\log p(\mathbf{x}) = -\frac{1}{2}(\mathbf{x}^\intercal \Lambda \mathbf{x} - 2\mathbf{x}^\intercal \Lambda \boldsymbol{\mu}) + K \tag{6}$$

where K is a constant that does not depend on $\mathbf{x}$.

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Claim 1.

$\rightarrow$)

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{D/2}\Lambda^{-\frac{1}{2}}} \exp\left\{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal \Lambda (\mathbf{x}-\boldsymbol{\mu})\right\}$$

$$\begin{aligned}
\log p(\mathbf{x}) &= -\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal \Lambda (\mathbf{x}-\boldsymbol{\mu}) - \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&= -\frac{1}{2}(\mathbf{x}^\intercal \Lambda \mathbf{x} - 2\mathbf{x}^\intercal \Lambda \boldsymbol{\mu}) - \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} - \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&= -\frac{1}{2}(\mathbf{x}^\intercal \Lambda \mathbf{x} - 2\mathbf{x}^\intercal \Lambda \boldsymbol{\mu}) + K
\end{aligned}$$

with $K = -\frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} - \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}})$.

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Claim 1.

$\leftarrow$)

$$
\begin{aligned}
\log p(\mathbf{x}) &= -\frac{1}{2}(\mathbf{x}^\intercal \Lambda \mathbf{x} - 2\mathbf{x}^\intercal \Lambda \boldsymbol{\mu}) + K \\
\log p(\mathbf{x}) &= -\frac{1}{2}(\mathbf{x}^\intercal \Lambda \mathbf{x} - 2\mathbf{x}^\intercal \Lambda \boldsymbol{\mu}) - \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} - \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&\quad + K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&= -\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal \Lambda (\mathbf{x}-\boldsymbol{\mu}) - \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&\quad + K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
&= \log N(\mathbf{x}|\boldsymbol{\mu}, \Lambda) + K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}}) \\
p(\mathbf{x}) &= N(\mathbf{x}|\boldsymbol{\mu}, \Lambda) \exp\left(K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}})\right)
\end{aligned}
\tag{7}
$$

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Claim 1.

$\leftarrow$) cont

$$
\begin{aligned}
1 &= \int p(\mathbf{x}) d\mathbf{x} \\
&= \int N(\mathbf{x}|\boldsymbol{\mu}, \Lambda) \exp\left(K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}})\right) d\mathbf{x} \\
&= \exp\left(K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}})\right) \int N(\mathbf{x}|\boldsymbol{\mu}, \Lambda) d\mathbf{x} \\
&= \exp\left(K + \frac{1}{2}\boldsymbol{\mu}^\intercal \Lambda \boldsymbol{\mu} + \log((2\pi)^{D/2}\Lambda^{-\frac{1}{2}})\right)
\end{aligned}
$$

From Eq. 7 then $p(\mathbf{x}) = N(\mathbf{x}|\boldsymbol{\mu}, \Lambda)$.

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Theorem 3, Eq. 1.

$$p(\mathbf{x}_a|\mathbf{x}_b) = \frac{p(\mathbf{x}_a, \mathbf{x}_b)}{p(\mathbf{x}_b)} = \frac{p(\mathbf{x})}{p(\mathbf{x}_b)}$$

$$\log p(\mathbf{x}_a|\mathbf{x}_b) = \log p(\mathbf{x}) - \log p(\mathbf{x}_b) = \log p(\mathbf{x}) + K$$

Therefore, the terms of $\log p(\mathbf{x}_a|\mathbf{x}_b)$ that depend on $\mathbf{x}_a$ are those of $\log p(\mathbf{x})$.
Steps for the proof:

1. isolate the terms of $\log p(\mathbf{x})$ that depend on $\mathbf{x}_a$,
2. notice that these term has the quadratic form of Claim 1, therefore $p(\mathbf{x}_a|\mathbf{x}_b)$ is Gaussian,
3. identify $\boldsymbol{\mu}$ and Λ in this quadratic form.

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 1)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Theorem 3, Eq. 1.

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{D/2}|\Lambda|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal \Lambda (\mathbf{x}-\boldsymbol{\mu})\right)$$

$$\begin{aligned}
\log p(\mathbf{x}) &= -\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal \Lambda (\mathbf{x}-\boldsymbol{\mu}) + K_1 \\
&= -\frac{1}{2}[(\mathbf{x}_a-\boldsymbol{\mu}_a)^\intercal, (\mathbf{x}_b-\boldsymbol{\mu}_b)^\intercal] \begin{bmatrix} \Lambda_{aa} & \Lambda_{ab} \\ \Lambda_{ba} & \Lambda_{bb} \end{bmatrix} \begin{bmatrix} \mathbf{x}_a-\boldsymbol{\mu}_a \\ \mathbf{x}_b-\boldsymbol{\mu}_b \end{bmatrix} + K_1 \\
&= -\frac{1}{2}\{(\mathbf{x}_a-\boldsymbol{\mu}_a)^\intercal \Lambda_{aa}(\mathbf{x}_a-\boldsymbol{\mu}_a) + 2(\mathbf{x}_a-\boldsymbol{\mu}_a)^\intercal \Lambda_{ab}(\mathbf{x}_b-\boldsymbol{\mu}_b) \\
&\quad + (\mathbf{x}_b-\boldsymbol{\mu}_b)^\intercal \Lambda_{bb}(\mathbf{x}_b-\boldsymbol{\mu}_b)\} + K_1 \\
&= -\frac{1}{2}\{\mathbf{x}_a^\intercal \Lambda_{aa}\mathbf{x}_a - 2\mathbf{x}_a^\intercal(\Lambda_{aa}\boldsymbol{\mu}_a - \Lambda_{ab}(\mathbf{x}_b-\boldsymbol{\mu}_b))\} + K_2 \\
&= -\frac{1}{2}\{\mathbf{x}_a^\intercal \Lambda_{aa}\mathbf{x}_a - 2\mathbf{x}_a^\intercal \Lambda_{aa}(\boldsymbol{\mu}_a - \Lambda_{aa}^{-1}\Lambda_{ab}(\mathbf{x}_b-\boldsymbol{\mu}_b))\} + K_2
\end{aligned}$$

Comparing the last equation with Eq. 6 we see that $\Lambda = \Lambda_{aa}$, $\boldsymbol{\mu} = \boldsymbol{\mu}_a - \Lambda_{aa}^{-1}\Lambda_{ab}(\mathbf{x}_b - \boldsymbol{\mu}_b)$ and conclude that

$$p(\mathbf{x}_a|\mathbf{x}_b) = \mathcal{N}(\mathbf{x}_a|\boldsymbol{\mu}_a - \Lambda_{aa}^{-1}\Lambda_{ab}(\mathbf{x}_b-\boldsymbol{\mu}_b), \Lambda_{aa})$$

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 2)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Claim 2 (Inverse of a partitioned matrix)

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix}^{-1} = \begin{pmatrix} M & -MBD^{-1} \\ -D^{-1}CM & D^{-1} + D^{-1}CMBD^{-1} \end{pmatrix} \quad (8)$$

where

$$M = (A - BD^{-1}C)^{-1}$$

Proof.

Exercise. Hint: verify that the multiplication of the inverse of the matrix in the right hand side of Eq. 8 with the matrix in the left hand side of the same equation is the identity matrix.

Proof: the conditional of a Gaussian is a Gaussian (Theorem 3, Eq. 2)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof of Theorem 3, Eq. 2.

Using the definition

$$\begin{pmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{pmatrix}^{-1} = \begin{pmatrix} \Lambda_{aa} & \Lambda_{ab} \\ \Lambda_{ba} & \Lambda_{bb} \end{pmatrix}$$

and using Eq. 8, we obtain

$$\Lambda_{aa} = (\Sigma_{aa} - \Sigma_{ab}\Sigma_{bb}^{-1}\Sigma_{ba})^{-1}$$
$$\Lambda_{ab} = -(\Sigma_{aa} - \Sigma_{ab}\Sigma_{bb}^{-1}\Sigma_{ba})^{-1}\Sigma_{ab}\Sigma_{bb}^{-1}$$

Replacing the above equations in Eq. 1 we obtain Eq. 2. □

Contents

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1 The Gaussian distribution

2 Linear models for regression
- Least-squares regression
- Maximum-likelihood regression
- Bayesian linear regression
  - Batch Bayesian linear regression
  - Online Bayesian linear regression

Linear regression example

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

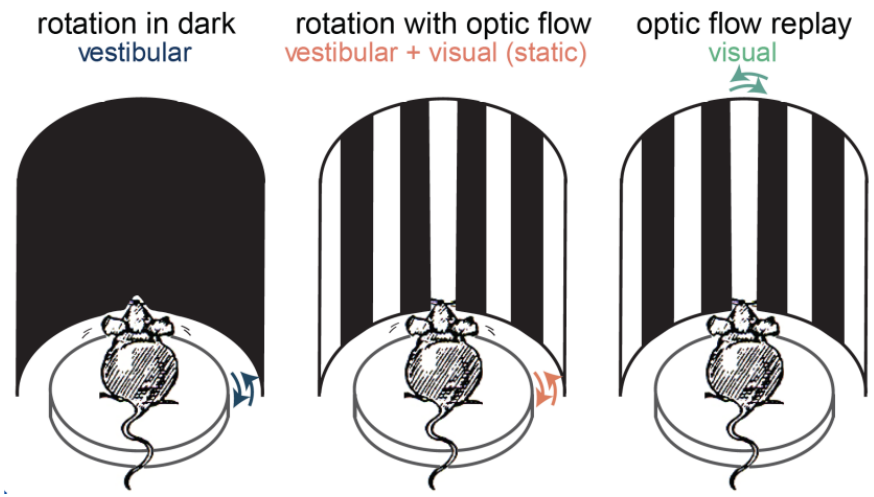

Keshavarzi et al., 2021

Linear regression example

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

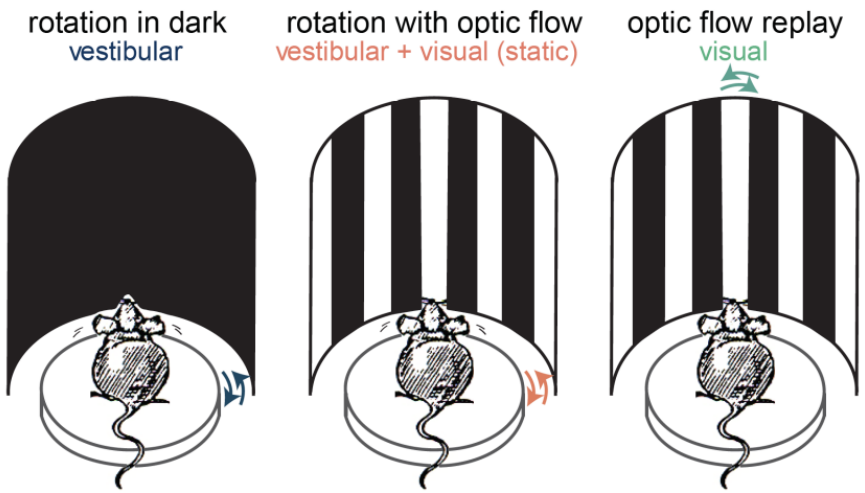

Keshavarzi et al., 2021

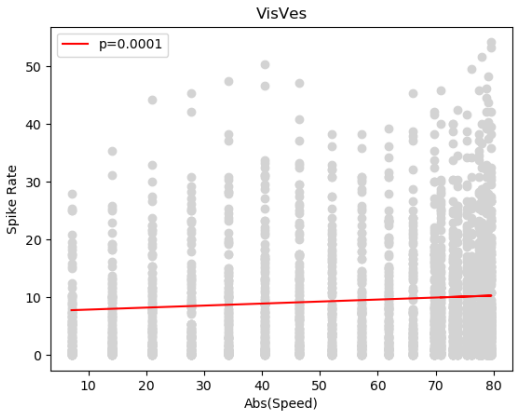

Linear regression example

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

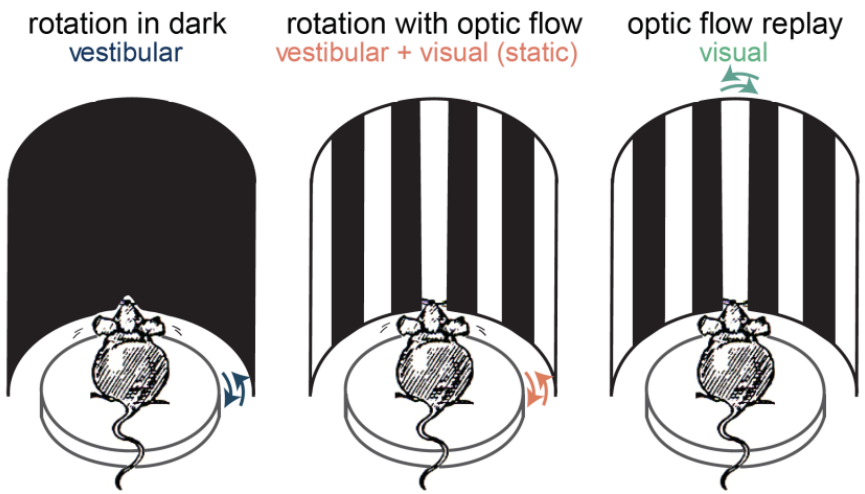

Keshavarzi et al., 2021

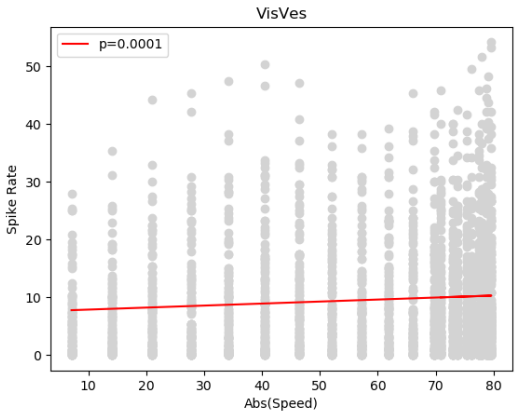

Is there a linear relation between the speed of rotation and the firing rate of visual cells?

Linear regression model

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

simple linear regression model

$$y(x_i, \mathbf{w}) = w_0 + w_1 x_i = [1, x_i] \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = [\phi_0(x_i), \phi_1(x_i)] \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$$
$$= \boldsymbol{\phi}(x_i)^\mathsf{T} \mathbf{w}$$

polynomial regression model

$$y(x_i, \mathbf{w}) = w_0 + w_1 x_i + w_2 x_i^2 + w_3 x_i^3 = [1, x_i, x_i^2, x_i^3] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix}$$
$$= [\phi_0(x_i), \phi_1(x_i), \phi_2(x_i), \phi_3(x_i)] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} = \boldsymbol{\phi}(x_i)^\mathsf{T} \mathbf{w}$$

basis functions linear regression model

$$y(x_i, \mathbf{w}) = \boldsymbol{\phi}(x_i)^\mathsf{T} \mathbf{w} = \sum_{j=1}^{M} w_j \phi_j(x_i)$$

Linear regression model

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

$$\mathbf{y}(\mathbf{x}, \mathbf{w}) = \begin{bmatrix} y(x_1, \mathbf{w}) \\ y(x_2, \mathbf{w}) \\ \dots \\ y(x_N, \mathbf{w}) \end{bmatrix} = \begin{bmatrix} \phi_1(x_1) & \phi_2(x_1) & \dots & \phi_M(x_1) \\ \phi_1(x_2) & \phi_2(x_2) & \dots & \phi_M(x_2) \\ \vdots & \vdots & \dots & \vdots \\ \phi_1(x_N) & \phi_2(x_N) & \dots & \phi_M(x_N) \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_M \end{bmatrix}$$

$$= \mathbf{\Phi w}$$

where $\mathbf{y}(\mathbf{x}, \mathbf{w}) \in \mathbb{R}^N$, $\mathbf{\Phi} \in \mathbb{R}^{N \times M}$, $\mathbf{w} \in \mathbb{R}^M$.

Basis functions for regression

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

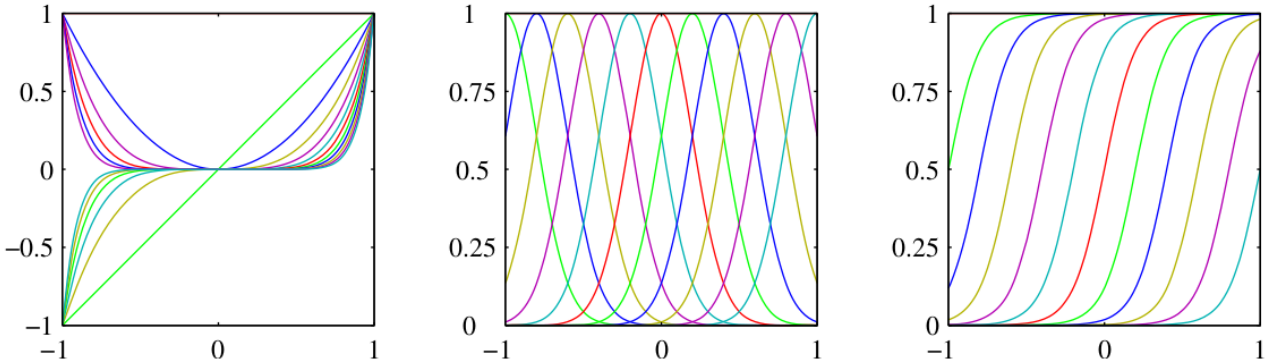

Figure 3.1 Examples of basis functions, showing polynomials on the left, Gaussians of the form (3.4) in the centre, and sigmoidal of the form (3.5) on the right.

Bishop (2016)

polynomial $\phi_i(x) = x^i$

Gaussian $\phi_i(x) = \exp(-\frac{(x-\mu_i)^2}{2\sigma^2})$

sigmoidal $\phi_i(x) = \frac{1}{1+\exp(-\frac{x-\mu_i}{\sigma^2})}$

Example dataset

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

Figure 1.2 Plot of a training data set of $N = 10$ points, shown as blue circles, each comprising an observation of the input variable x along with the corresponding target variable t. The green curve shows the function $\sin(2\pi x)$ used to generate the data. Our goal is to predict the value of t for some new value of x, without knowledge of the green curve.

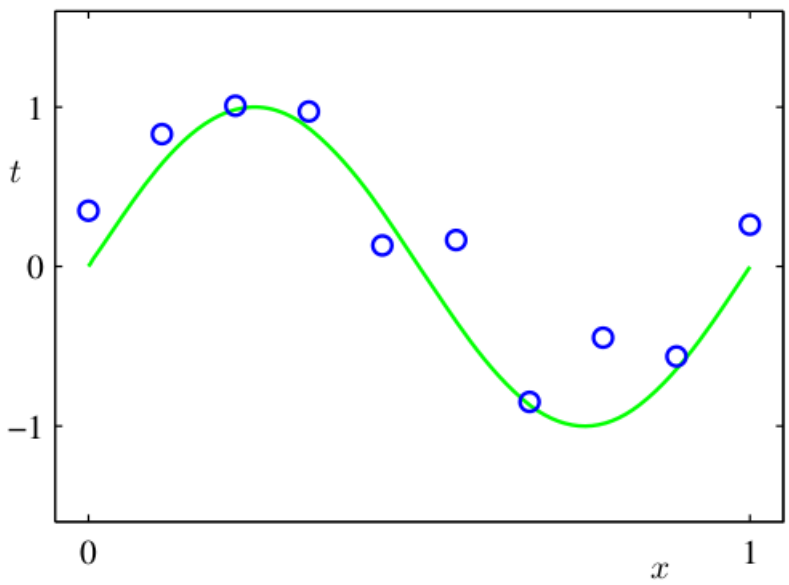

Bishop (2016)

Outline

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1 The Gaussian distribution

2 Linear models for regression
- Least-squares regression
- Maximum-likelihood regression
- Bayesian linear regression
  - Batch Bayesian linear regression
  - Online Bayesian linear regression

Least-squares estimation of model parameters (Trefethen and Bau III, 1997)

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Definition 2 (Least-squares problem)

Given $\mathbf{\Phi} \in \mathbb{R}^{N \times M}, N \geq M, \mathbf{t} \in \mathbb{R}^N$, *find* $\mathbf{w} \in \mathbb{R}^M$ *such that* $E_{LS}(\mathbf{w}) = ||\mathbf{t} - \mathbf{\Phi w}||_2$ *is minimised.*

Theorem 5 (Least-squares solution)

Let $\mathbf{\Phi} \in \mathbb{R}^{N \times M} (N \geq M)$ *and* $\mathbf{t} \in \mathbb{R}^N$ *be given. A vector* $\mathbf{w} \in \mathbb{R}^M$ *minimises* $||\mathbf{r}||_2 = ||\mathbf{t} - \mathbf{\Phi w}||_2$, *thereby solving the least-squares problem, if and only if* $\mathbf{r} \perp$ *range*$(\mathbf{\Phi})$, *that is,* $\mathbf{\Phi}^\mathsf{T}\mathbf{r} = 0$, *or equivalently,* $\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi w} = \mathbf{\Phi}^\mathsf{T}\mathbf{t}$, *or again equivalently,* $P\mathbf{t} = \mathbf{\Phi w}$.

Figure 3.2 Geometrical interpretation of the least-squares solution, in an N-dimensional space whose axes are the values of $t_1, \ldots, t_N$. The least-squares regression function is obtained by finding the orthogonal projection of the data vector $\mathbf{t}$ onto the subspace spanned by the basis functions $\phi_j(\mathbf{x})$ in which each basis function is viewed as a vector φ_j of length N with elements $\phi_j(\mathbf{x}_n)$.

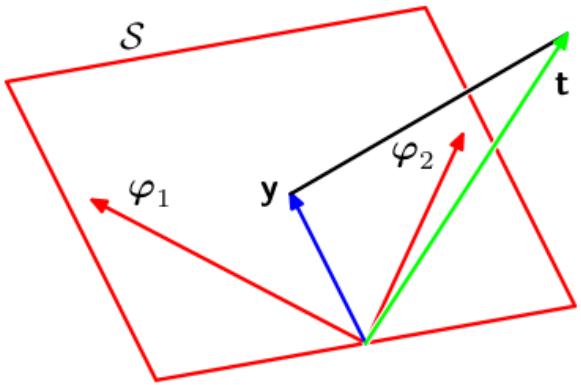

Bishop (2016)

Code for least-squares estimation of model parameters

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

- overfitting
- cross validation
- larger datasets allow more complex models

Regularised least-squares estimation of model parameters

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

To cope with the overfitting of least squares, we can add to the least squares optimisation criterion a term that enforces coefficients to be zero. The regularised least-squares optimisation criterion becomes:

$$E_{RLS}(\mathbf{w}) = ||\mathbf{t} - \mathbf{\Phi w}||_2^2 + \lambda ||\mathbf{w}||_2^2$$

where λ is the regularisation parameter that weights the strength of the regularisation.

Regularised least-squares estimation of model parameters

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Claim 3 (Regularized least-squares estimate)

$$\mathbf{w}_{RLS} = \arg\min_{\mathbf{w}} E_{RLS}(\mathbf{w}) = \arg\min_{\mathbf{w}} ||\mathbf{t} - \mathbf{\Phi w}||_2^2 + \lambda ||\mathbf{w}||_2^2 = (\lambda \mathbf{I} + \mathbf{\Phi}^\intercal \mathbf{\Phi})^{-1} \mathbf{\Phi}^\intercal \mathbf{t}$$

Proof.

Since $E_{RLS}(\mathbf{w})$ is a polynomial of order two on the elements of $\mathbf{w}$ (i.e., a quadratic form), we can use the *Completing the Squares* technique below to find its minimum.

$$\begin{aligned}
\boldsymbol{\mu} &= \arg\max_{\mathbf{w}} \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}, \Sigma) = \arg\max_{\mathbf{w}} \log \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}, \Sigma) \\
&= \arg\max_{\mathbf{w}} \{K - \frac{1}{2}(-2\boldsymbol{\mu}^\intercal \Sigma^{-1}\mathbf{w} + \mathbf{w}\Sigma^{-1}\mathbf{w})\} \qquad (9) \\
&= \arg\min_{\mathbf{w}} \{-K + \frac{1}{2}(-2\boldsymbol{\mu}^\intercal \Sigma^{-1}\mathbf{w} + \mathbf{w}\Sigma^{-1}\mathbf{w})\} \\
&= \arg\min_{\mathbf{w}} \{K_1 - 2\boldsymbol{\mu}^\intercal \Sigma^{-1}\mathbf{w} + \mathbf{w}\Sigma^{-1}\mathbf{w}\} \qquad (10)
\end{aligned}$$

Note: Eq. 9 uses Eq. 6.
To find the minimum of a quadratic form, we write it in the form of the terms inside the curly brackets of Eq. 10, and the term corresponding to $\boldsymbol{\mu}$ will be the minimum.

Regularised least-squares estimation of model parameters

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof.

Let's write E_{RLS} in the form of the terms inside the curly brackets of Eq. 10.

$$\begin{aligned} E_{RLS} &= ||\mathbf{t} - \mathbf{\Phi w}||_2^2 + \lambda ||\mathbf{w}||_2^2 = (\mathbf{t} - \mathbf{\Phi w})^\mathsf{T}(\mathbf{t} - \mathbf{\Phi w}) + \lambda \mathbf{w}^\mathsf{T}\mathbf{w} \\ &= \mathbf{t}^\mathsf{T}\mathbf{t} - 2\mathbf{t}^\mathsf{T}\mathbf{\Phi w} + \mathbf{w}^\mathsf{T}\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi w} + \lambda \mathbf{w}^\mathsf{T}\mathbf{w} \\ &= \mathbf{t}^\mathsf{T}\mathbf{t} - 2\mathbf{t}^\mathsf{T}\mathbf{\Phi w} + \mathbf{w}^\mathsf{T}(\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi}^\mathsf{T} + \lambda \mathbf{I}_M)\mathbf{w} \end{aligned}$$

Calling

$$\begin{aligned} \Sigma^{-1} &= \mathbf{\Phi}^\mathsf{T}\mathbf{\Phi}^\mathsf{T} + \lambda \mathbf{I}_M \\ \boldsymbol{\mu}^\mathsf{T}\Sigma^{-1} &= \mathbf{t}^\mathsf{T}\mathbf{\Phi} \text{ or } \boldsymbol{\mu}^\mathsf{T} = \mathbf{t}^\mathsf{T}\mathbf{\Phi}\Sigma \text{ or } \boldsymbol{\mu} = \Sigma\mathbf{\Phi}^\mathsf{T}\mathbf{t} = \left(\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi}^\mathsf{T} + \lambda \mathbf{I}_M\right)^{-1}\mathbf{\Phi}^\mathsf{T}\mathbf{t} \end{aligned}$$

we can express

$$E_{RLS} = K + 2\boldsymbol{\mu}^\mathsf{T}\Sigma^{-1}\mathbf{w} + \mathbf{w}\Sigma^{-1}\mathbf{w}$$

Then

$$\mathbf{w}_{RLS} = \arg\min_{\mathbf{w}} E_{RLS}(\mathbf{w}) = \boldsymbol{\mu} = \left(\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi}^\mathsf{T} + \lambda \mathbf{I}_M\right)^{-1}\mathbf{\Phi}^\mathsf{T}\mathbf{t}$$

□

Code for regularised least-squares estimation of model parameters

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

- control of overfitting

Outline

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1 The Gaussian distribution

2 Linear models for regression

- Least-squares regression
- Maximum-likelihood regression
- Bayesian linear regression
  - Batch Bayesian linear regression
  - Online Bayesian linear regression

Maximum-likelihood estimation of model parameters

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

Definition 3 (Likelihood function)

For a statistical model characterised by a probability density function $p(\mathbf{x}|\theta)$ (or probability mass function $P_\theta(X = \mathbf{x})$) the likelihood function is a function of the parameters θ, $\mathcal{L}(\theta) = p(\mathbf{x}|\theta)$ (or $\mathcal{L}(\theta) = P_\theta(\mathbf{x})$).

Definition 4 (Maximum likelihood parameters estimates)

The maximum likelihood parameters estimates are the parameters that maximise the likelihood function.

$$\theta_{ML} = \arg\max_\theta \mathcal{L}(\theta)$$

Maximum-likelihood estimation for the basis function linear regression model

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

We seek the parameter $\mathbf{w}_{ML}$ and β_{ML} that maximised the following likelihood function

$$\mathcal{L}(\mathbf{w}, \beta) = p(\mathbf{t}|\mathbf{w}, \beta) = \mathcal{N}(\mathbf{t}|\mathbf{\Phi}\mathbf{w}, \beta^{-1} I_N) = \prod_{n=1}^{N} \mathcal{N}(t_n|\phi^{\mathsf{T}}(x_n)\mathbf{w}, \beta^{-1}) \quad (11)$$

They are

$$\mathbf{w}_{ML} = (\mathbf{\Phi}^{\mathsf{T}}\mathbf{\Phi})^{-1}\mathbf{\Phi}^{\mathsf{T}}\mathbf{t} \quad (12)$$

$$\frac{1}{\beta_{ML}} = \frac{1}{N}\sum_{n=1}^{N}(t_n - \phi(\mathbf{x}_n)^{\mathsf{T}}\mathbf{w}_{ML})^2 \quad (13)$$

- first regression method that assumes random observations
- if the likelihood function is assumed to be Normal, maximum-likelihood and least-squares coefficients estimates are equal.

Exercise

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

Exercise 1

Derive the formulas for the maximum likelihood estimates of the coefficients, $\mathbf{w}$*, and noise precision,* β*, of the basis functions linear regression model given in Eqs. 12 and 13.*

Outline

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

1 The Gaussian distribution

2 Linear models for regression
- Least-squares regression
- Maximum-likelihood regression
- Bayesian linear regression
  - Batch Bayesian linear regression
  - Online Bayesian linear regression

Bayesian linear regression: motivation

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

- elegant,
- naturally allows online regression,
- does not require cross-validation for model selection,
- it is the first step to more complex Bayesian modelling.

Batch Bayesian linear regression: posterior distribution of parameters

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

In Bayesian linear regression we seek the posterior distribution of the weights of the linear regression model, $\mathbf{w}$, given the observations, which is proportional to the product of the likelihood function, $p(\mathbf{t}|\mathbf{w})$, and the prior, $p(\mathbf{w})$; i.e.,

$$p(\mathbf{w}|\mathbf{t}) \propto p(\mathbf{t}|\mathbf{w})p(\mathbf{w})$$

To calculate this posterior below we use the likelihood function defined in Eq. 11 and the following prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$$

Using the expression of the conditional of the Linear Gaussian model, Eq. 5, we obtain

$$p(\mathbf{w}|\mathbf{t}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

$$\mathbf{m}_N = \beta \mathbf{S}_N \mathbf{\Phi}^\mathsf{T} \mathbf{t} \tag{14}$$

$$\mathbf{S}_N^{-1} = \alpha\mathbf{I} + \beta\mathbf{\Phi}^\mathsf{T}\mathbf{\Phi} \tag{15}$$

Batch Bayesian linear regression: exercise

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Exercise 2

Derive the formulas for the Bayesian posterior mean (Eq. 14) and covariance (Eq. 15) of the basis function linear regression model.

Exercise 3

Show that

$$log \log p(\mathbf{w}|\boldsymbol{t}) = K - \frac{\beta}{2}||\mathbf{t} - \mathbf{\Phi w}||_2^2 - \frac{\alpha}{2}||\mathbf{w}||_2^2 \tag{16}$$

Therefore, the maximum-a-posteriori parameters of the basis function linear regression model are the solution of the regularized least-squares problem with $\lambda = \alpha/\beta$.
Note that, as we will show next, Bayesian linear regression uses the full posterior of the parameters to make predictions or to do model selection, and not just the maximum-a-posteriori parameters.

Batch Bayesian linear regression: demo code

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

Available here

Online Bayesian linear regression: recursive update of posterior distribution of parameters

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Claim 4 (recursive update)

If the observations, $\{\mathbf{t}_1, \ldots, \mathbf{t}_n, \ldots\}$, are linearly independent when conditioned on the model parameters, $\boldsymbol{\theta}$, then for any $n \in \mathbb{N}$

$$p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_n) = K\ p(\mathbf{t}_n|\boldsymbol{\theta}) p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_{n-1}) \tag{17}$$

where K is a quantity that does not depend on $\boldsymbol{\theta}$.

Online Bayesian linear regression: recursive update of posterior distribution of parameters

The Gaussian distribution
Linear models for regression
Least-squares regression
Maximum-likelihood regression
Bayesian linear regression
Batch Bayesian linear regression
Online Bayesian linear regression
References

Proof.

By induction on $H_n : p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_n) = K\ p(\mathbf{t}_n|\boldsymbol{\theta})p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_{n-1})$.

H_1

$$p(\boldsymbol{\theta}|\mathbf{t}_1) = \frac{p(\boldsymbol{\theta}, \mathbf{t}_1)}{p(\mathbf{t}_1)} = \frac{p(\mathbf{t}_1|\boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathbf{t}_1)} = K\ p(\mathbf{t}_1|\boldsymbol{\theta})p(\boldsymbol{\theta})$$

$H_n \rightarrow H_{n+1}$

$$\begin{aligned}
p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_{n+1}) &= \frac{p(\boldsymbol{\theta}, \mathbf{t}_1, \ldots, \mathbf{t}_{n+1})}{p(\mathbf{t}_1, \ldots, \mathbf{t}_{n+1})} \\
&= \frac{p(\mathbf{t}_{n+1}|\boldsymbol{\theta}, \mathbf{t}_1, \ldots, \mathbf{t}_n)p(\boldsymbol{\theta}, \mathbf{t}_1, \ldots, \mathbf{t}_n)}{p(\mathbf{t}_1 \ldots, \mathbf{t}_{n+1})} \\
&= \frac{p(\mathbf{t}_{n+1}|\boldsymbol{\theta})p(\boldsymbol{\theta}, \mathbf{t}_1, \ldots, \mathbf{t}_n)}{p(\mathbf{t}_1 \ldots, \mathbf{t}_{n+1})} \\
&= \frac{p(\mathbf{t}_{n+1}|\boldsymbol{\theta})p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_n)p(\mathbf{t}_1, \ldots, \mathbf{t}_n)}{p(\mathbf{t}_1 \ldots, \mathbf{t}_{n+1})} \\
&= K\ p(\mathbf{t}_{n+1}|\boldsymbol{\theta})p(\boldsymbol{\theta}|\mathbf{t}_1, \ldots, \mathbf{t}_n)
\end{aligned}$$

Note: the third equality above holds because the observations are assumed to be conditional independent given the parameters.

□

References

The Gaussian distribution

Linear models for regression

Least-squares regression

Maximum-likelihood regression

Bayesian linear regression

Batch Bayesian linear regression

Online Bayesian linear regression

References

Bishop, C. M. (2016). *Pattern recognition and machine learning*. Springer-Verlag New York.

Cover, T. M. and Thomas, J. A. (1991). *Elements of information theory*. John Wiley & Sons.

Durbin, J. and Koopman, S. J. (2012). *Time series analysis by state space methods*, volume 38. OUP Oxford.

Papoulis, A. and Pillai, S. U. (2002). *Probability, random variables and stochastic processes*. Mc Graw Hill, fourth edition.

Trefethen, L. n. and Bau III, D. (1997). Numerical linear algebra.

Williams, C. K. and Rasmussen, C. E. (2006). *Gaussian processes for machine learning*, volume 2. MIT press Cambridge, MA.

Yu, B. M., Cunningham, J. P., Santhanam, G., Ryu, S. I., Shenoy, K. V., and Sahani, M. (2009). Gaussian-process factor analysis for low-dimensional single-trial analysis of neural population activity. *Journal of neurophysiology*, 102(1):614–635.